

# PIG USE CASE: COGNIZANT EMPLOYEE DATA ANALYSIS



BIG DATA  
HADOOP  
&  
SPARK  
TRAINING

ACADGILD

ASSIGNMENT  
5.1

BY :-

SAHIL  
KHURANA

## Problem Statement

- (a) Top 5 employees (employee id and employee name) with highest rating. (In case two employees have same rating, employee with name coming first in dictionary should get preference)
- (b) Top 3 employees (employee id and employee name) with highest salary, whose employee id is an odd number. (In case two employees have same salary, employee with name coming first in dictionary should get preference)
- (c) Employee (employee id and employee name) with maximum expense (In case two employees have same expense, employee with name coming first in dictionary should get preference)
- (d) List of employees (employee id and employee name) having entries in employee\_expenses file.
- (e) List of employees (employee id and employee name) having no entry in employee\_expenses file.

## Associated Data Files

We have employee\_details and employee\_expenses files. Use local mode while running Pig and write Pig Latin script to get below results:

employee\_details.txt

[https://github.com/prateekATacadgild/DatasetsForCognizant/blob/master/employee\\_details.txt](https://github.com/prateekATacadgild/DatasetsForCognizant/blob/master/employee_details.txt)

```
101,Amitabh,20000,1
102,Shahrukh,10000,2
103,Akshay,11000,3
104,Anubhav,5000,4
105,Pawan,2500,5
106,Aamir,25000,1
107,Salman,17500,2
108,Ranbir,14000,3
109,Katrina,1000,4
110,Priyanka,2000,5
111,Tushar,500,1
112,Ajay,5000,2
113,Jubeen,1000,1
114,Madhuri,2000,2
```

employee\_expenses.txt

[https://github.com/prateekATacadgild/DatasetsForCognizant/blob/master/employee\\_expenses.txt](https://github.com/prateekATacadgild/DatasetsForCognizant/blob/master/employee_expenses.txt)

```
101 200
102 100
```

110 400  
114 200  
119 200  
105 100  
101 100  
104 300  
102 400

Note: - To solve the Assignment, I have created a VM with Ubuntu 16.04 OS and configured Hadoop 2.8.2 and pig-0.17.0 on the same

Put the dataset in HDFS location

```
sahil@ubuntu:~/Desktop$ hdfs dfs -mkdir /u01/hadoop/Pig/cognizant_employee_data_usecase
sahil@ubuntu:~/Desktop$ hdfs dfs -put employee_details.txt /u01/hadoop/Pig/cognizant_employee_data_usecase
sahil@ubuntu:~/Desktop$ hdfs dfs -put employee_expenses.txt /u01/hadoop/Pig/cognizant_employee_data_usecase
sahil@ubuntu:~/Desktop$ hdfs dfs -ls /u01/hadoop/Pig/cognizant_employee_data_usecase
Found 2 items
-rw-r--r-- 1 sahil supergroup 261 2017-12-05 23:11 /u01/hadoop/Pig/cognizant_employee_data_usecase/employee_details.txt
-rw-r--r-- 1 sahil supergroup 72 2017-12-05 23:12 /u01/hadoop/Pig/cognizant_employee_data_usecase/employee_expenses.txt
sahil@ubuntu:~/Desktop$
```

## Load dataset in Pig

### employee\_details.txt

```
grunt> employee_details = LOAD
```

```
'/u01/hadoop/Pig/cognizant_employee_data_usecase/employee_details.txt' USING
PigStorage(',') AS (id:int, name:chararray, salary:int, ratings:int);
```

```
grunt> describe employee_details;
```

```
sahil@ubuntu:~/Desktop$ pig -x mapreduce
17/12/06 00:57:18 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL
17/12/06 00:57:18 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE
17/12/06 00:57:18 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType
2017-12-06 00:57:18,362 [main] INFO org.apache.pig.Main - Apache Pig version 0.17.0 (r1797386) compiled Jun 02 2017, 15:41:58
2017-12-06 00:57:18,363 [main] INFO org.apache.pig.Main - Logging error messages to: /home/sahil/Desktop/pig_1512550638361.log
2017-12-06 00:57:18,592 [main] INFO org.apache.pig.impl.util.Utils - Default bootup file /home/sahil/.pigbootup not found
2017-12-06 00:57:22,756 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2017-12-06 00:57:22,756 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to hadoop file system at: hdfs://localhost:9000
2017-12-06 00:57:27,668 [main] INFO org.apache.pig.PigServer - Pig Script ID for the session: PIG-default-4417480d-40e5-4d85-93bf-74e2d7b9473a
2017-12-06 00:57:27,668 [main] WARN org.apache.pig.PigServer - ATS is disabled since yarn.timeline-service.enabled set to false
grunt> employee_details = LOAD '/u01/hadoop/Pig/cognizant_employee_data_usecase/employee_details.txt' USING PigStorage(',') AS (id:int, name:chararray, salary:int, ratings:int);
grunt> describe employee_details;
employee_details: {id: int,name: chararray,salary: int,ratings: int}
grunt>
```

```
grunt> dump employee_details;
```

```
sahil@ubuntu: ~/Desktop
us=SUCCEEDED. Redirecting to
2017-12-06 01:23:11,119 [ma
2017-12-06 01:23:11,124 [ma
us=SUCCEEDED. Redirecting to
2017-12-06 01:23:11,163 [ma
2017-12-06 01:23:11,166 [ma
2017-12-06 01:23:11,192 [ma
2017-12-06 01:23:11,194 [ma
(101,Amitabh,20000,1)
(102,Shahrukh,10000,2)
(103,Akshay,11000,3)
(104,Anubhav,5000,4)
(105,Pawan,2500,5)
(106,Aamir,25000,1)
(107,Salman,17500,2)
(108,Ranbir,14000,3)
(109,Katrina,1000,4)
(110,Priyanka,2000,5)
(111,Tushar,500,1)
(112,Ajay,5000,2)
(113,Jubeen,1000,1)
(114,Madhuri,2000,2)
grunt>
grunt>
```

## employee\_expenses.txt

```
grunt> employee_expenses = LOAD
'/u01/hadoop/Pig/cognizant_employee_data_usecase/employee_expenses.txt' USING
PigStorage('\t') AS (id:int, expenses:int);
```

```
grunt>
grunt> employee_expenses = LOAD '/u01/hadoop/Pig/cognizant_employee_data_usecase/employee_expenses.txt' USING PigStorage('\t') AS (id:int,
expenses:int);
grunt> dump employee_expenses;
```

```
grunt> dump employee_expenses;
```

```
sahil@ubuntu: ~/Desktop
2017-12-06 02:22:17,619 [main]
2017-12-06 02:22:17,692 [main]
2017-12-06 02:22:17,692 [main]
(101,200)
(102,100)
(110,400)
(114,200)
(119,200)
(105,100)
(101,100)
(104,300)
(102,400)
grunt>
```

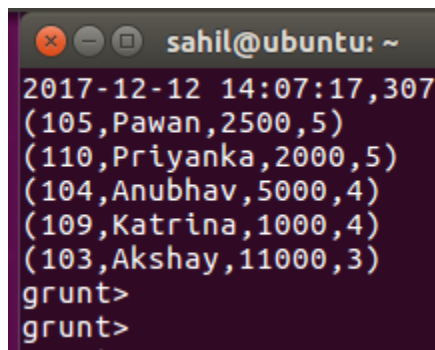
**(a) Top 5 employees (employee id and employee name) with highest rating. (In case two employees have same rating, employee with name coming first in dictionary should get preference)**

Commands Used:-

```
grunt> employee_details_order = order employee_details by ratings DESC, name;
grunt> employee_details_order_limit_5 = limit employee_details_order 5;
grunt> DUMP employee_details_order_limit_5;
```

```
grunt>
grunt> employee_details_order = order employee_details by ratings DESC, name;
grunt> employee_details_order_limit_5 = limit employee_details_order 5;
grunt> DUMP employee_details_order_limit_5; █
```

Result:-



```
sahil@ubuntu: ~
2017-12-12 14:07:17,307
(105,Pawan,2500,5)
(110,Priyanka,2000,5)
(104,Anubhav,5000,4)
(109,Katrina,1000,4)
(103,Akshay,11000,3)
grunt>
grunt>
```

(105,Pawan,2500,5)  
 (110,Priyanka,2000,5)  
 (104,Anubhav,5000,4)  
 (109,Katrina,1000,4)  
 (103,Akshay,11000,3)

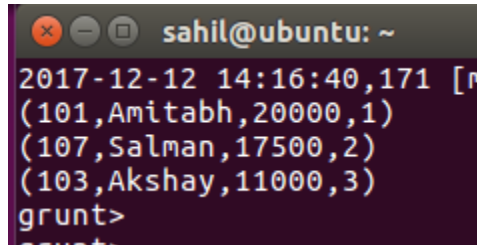
**(b) Top 3 employees (employee id and employee name) with highest salary, whose employee id is an odd number. (In case two employees have same salary, employee with name coming first in dictionary should get preference)**

Commands Used:-

```
grunt> employee_details_filter = FILTER employee_details by (id%2!=0);
grunt> employee_details_order = ORDER employee_details_filter by salary DESC,
name;
grunt> employee_details_limit = LIMIT employee_details_order 3;
grunt> DUMP employee_details_limit;
```

```
grunt>
grunt> employee_details_filter = FILTER employee_details by (id%2!=0);
grunt> employee_details_order = ORDER employee_details_filter by salary DESC, name;
grunt> employee_details_limit = LIMIT employee_details_order 3;
grunt> DUMP employee_details_limit;
```

Result:-



```
2017-12-12 14:16:40,171 [r
(101,Amitabh,20000,1)
(107,Salman,17500,2)
(103,Akshay,11000,3)
grunt>
```

**(101,Amitabh,20000,1)**

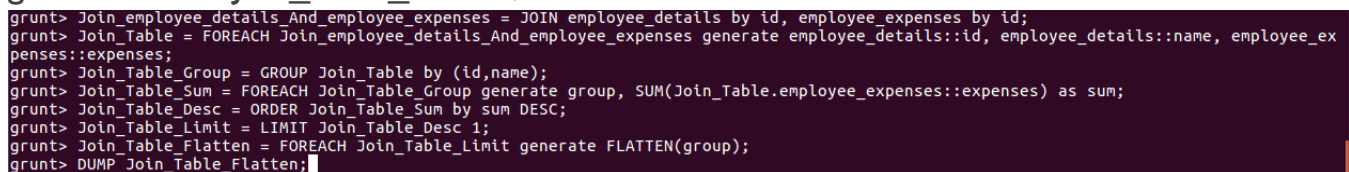
**(107,Salman,17500,2)**

**(103,Akshay,11000,3)**

**(c) Employee (employee id and employee name) with maximum expense (In case two employees have same expense, employee with name coming first in dictionary should get preference)**

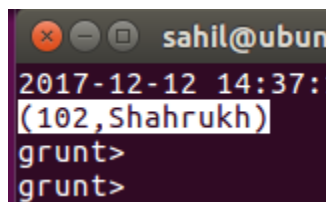
Commands Used:-

```
grunt> Join_employee_details_And_employee_expenses = JOIN employee_details by id,
employee_expenses by id;
grunt> Join_Table = FOREACH Join_employee_details_And_employee_expenses
generate employee_details::id, employee_details::name, employee_expenses::expenses;
grunt> Join_Table_Group = GROUP Join_Table by (id,name);
grunt> Join_Table_Sum = FOREACH Join_Table_Group generate group,
SUM(Join_Table.employee_expenses::expenses) as sum;
grunt> Join_Table_Desc = ORDER Join_Table_Sum by sum DESC;
grunt> Join_Table_Limit = LIMIT Join_Table_Desc 1;
grunt> Join_Table_Flatten = FOREACH Join_Table_Limit generate FLATTEN(group);
grunt> DUMP Join_Table_Flatten;
```



```
grunt> Join_employee_details_And_employee_expenses = JOIN employee_details by id, employee_expenses by id;
grunt> Join_Table = FOREACH Join_employee_details_And_employee_expenses generate employee_details::id, employee_details::name, employee_ex
penses::expenses;
grunt> Join_Table_Group = GROUP Join_Table by (id,name);
grunt> Join_Table_Sum = FOREACH Join_Table_Group generate group, SUM(Join_Table.employee_expenses::expenses) as sum;
grunt> Join_Table_Desc = ORDER Join_Table_Sum by sum DESC;
grunt> Join_Table_Limit = LIMIT Join_Table_Desc 1;
grunt> Join_Table_Flatten = FOREACH Join_Table_Limit generate FLATTEN(group);
grunt> DUMP Join_Table_Flatten;
```

Result:-



```
2017-12-12 14:37:
(102,Shahrukh)
grunt>
grunt>
```

**(102,Shahrukh)**

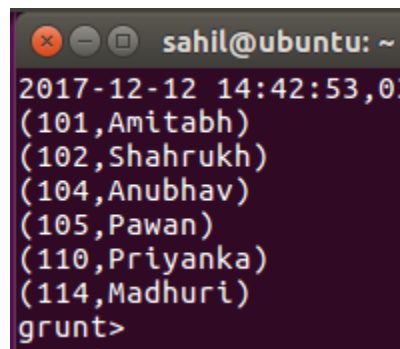
### (d) List of employees (employee id and employee name) having entries in employee\_expenses file.

Commands Used:-

```
grunt> Join_employee_details_And_employee_expenses = JOIN employee_details by id,
employee_expenses by id;
grunt> Join_Table_2 = FOREACH Join_employee_details_And_employee_expenses
generate employee_details::id, employee_details::name;
grunt> Join_Table_2_Distinct = DISTINCT Join_Table_2;
grunt> dump Join_Table_2_Distinct;
```

```
grunt>
grunt> Join_Table_2 = FOREACH Join_employee_details_And_employee_expenses generate employee_details::id, employee_details::name;
grunt> Join_Table_2_Distinct = DISTINCT Join_Table_2;
grunt> dump Join_Table_2_Distinct;
```

Result:-



```
sahil@ubuntu: ~
2017-12-12 14:42:53,01
(101,Amitabh)
(102,Shahrukh)
(104,Anubhav)
(105,Pawan)
(110,Priyanka)
(114,Madhuri)
grunt>
```

(101,Amitabh)  
 (102,Shahrukh)  
 (104,Anubhav)  
 (105,Pawan)  
 (110,Priyanka)  
 (114,Madhuri)

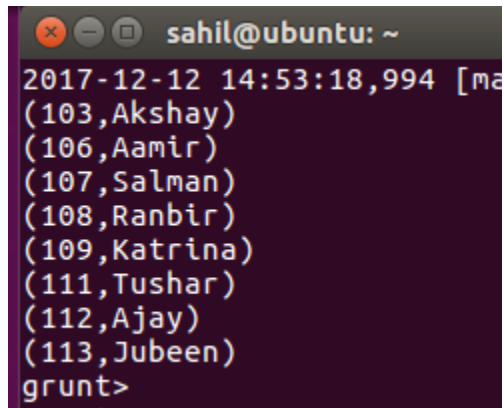
### (e) List of employees (employee id and employee name) having no entry in employee\_expenses file.

Commands Used:-

```
grunt> Left_Outer_Join_employee_details_And_employee_expenses = JOIN
employee_details by id LEFT OUTER, employee_expenses by id;
grunt> Left_Outer_Join_Filter = FILTER
Left_Outer_Join_employee_details_And_employee_expenses by
employee_expenses::expenses is NULL;
grunt> Join_Table_3 = FOREACH Left_Outer_Join_Filter generate employee_details::id,
employee_details::name;
grunt> DUMP Join_Table_3;
```

```
grunt>
grunt> Left_Outer_Join_employee_details_And_employee_expenses = JOIN employee_details by id LEFT OUTER, employee_expenses by id;
grunt> Left_Outer_Join_Filter = FILTER Left_Outer_Join_employee_details_And_employee_expenses by employee_expenses::expenses is NULL;
grunt> Join_Table_3 = FOREACH Left_Outer_Join_Filter generate employee_details::id, employee_details::name;
grunt> DUMP Join_Table_3;
```

Result:-

A terminal window titled 'sahil@ubuntu: ~' with a dark background. It displays a list of names and IDs in a light blue font. The text is as follows:

```
2017-12-12 14:53:18,994 [ma
(103,Akshay)
(106,Aamir)
(107,Salman)
(108,Ranbir)
(109,Katrina)
(111,Tushar)
(112,Ajay)
(113,Jubeen)
grunt>
```

**(103,Akshay)**  
**(106,Aamir)**  
**(107,Salman)**  
**(108,Ranbir)**  
**(109,Katrina)**  
**(111,Tushar)**  
**(112,Ajay)**  
**(113,Jubeen)**